

Financial Risk Analysis with Python (Goldman Sachs Financial Analysis)

```
In [2]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sns
```

```
In [3]: goldman_df = pd.read_csv("goldman_sachs.csv")
print(goldman_df)
```

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	
0	33	CUST6549	ACC12334	Credit	Withdrawal	
1	177	CUST2942	ACC52650	Credit	Withdrawal	
2	178	CUST6776	ACC45101	Current	Deposit	
3	173	CUST2539	ACC88252	Current	Withdrawal	
4	67	CUST2626	ACC21878	Savings	Withdrawal	
..
795	11	CUST8461	ACC60432	Current	Withdrawal	
796	44	CUST1121	ACC41829	Credit	Withdrawal	
797	160	CUST3059	ACC28292	Current	Payment	
798	37	CUST1042	ACC28295	Current	Payment	
799	101	CUST2464	ACC53865	Loan	Deposit	
	Product	Firm	Region	Manager	TransactionDate	\
0	Savings Account	Firm C	Central	Manager 1	21-10-2023	
1	Home Loan	Firm A	East	Manager 3	20-06-2023	
2	Personal Loan	Firm C	South	Manager 3	02-01-2023	
3	Mutual Fund	Firm A	Central	Manager 2	25-07-2023	
4	Home Loan	Firm C	Central	Manager 4	25-07-2023	
..
795	Credit Card	Firm E	Central	Manager 3	17-02-2023	
796	Personal Loan	Firm C	East	Manager 4	26-10-2023	
797	Mutual Fund	Firm E	South	Manager 4	11-05-2023	
798	Personal Loan	Firm E	West	Manager 3	26-03-2023	
799	Mutual Fund	Firm C	South	Manager 4	19-07-2023	
	TransactionAmount	AccountBalance	RiskScore	CreditRating	TenureMonths	
0	87480.05448	74008.43310	0.729101	319	200	
1	20315.74505	22715.83590	0.472424	692	47	
2	10484.57165	42706.09210	0.648784	543	109	
3	45122.27373	114176.56870	0.734832	430	103	
4	42360.79878	17863.02644	0.289304	468	234	
..
795	119878.30500	72041.03969	0.437613	772	217	
796	97933.33752	98317.05068	0.240792	411	78	
797	63072.38174	35088.75351	0.283858	486	215	
798	-25289.82472	40794.07619	0.515020	549	238	
799	23936.04314	59787.07004	0.474789	656	164	

[800 rows x 15 columns]

```
In [4]: goldman_df.describe()
```

	TransactionID	TransactionAmount	AccountBalance	RiskScore	CreditRating	TenureMonths
count	800.000000	800.000000	800.000000	800.000000	800.000000	800.000000
mean	98.401250	51575.671765	72345.657910	0.473415	566.586250	126.58250
std	56.292618	29055.838886	34070.038539	0.242858	151.674261	68.98532
min	1.000000	-30721.247890	-37293.600250	-0.389354	305.000000	6.00000
25%	50.000000	31692.004800	49157.502723	0.313286	436.500000	65.00000
50%	97.000000	50249.069385	72789.370195	0.472155	563.000000	132.00000
75%	146.000000	71913.394570	95352.455698	0.636848	696.000000	186.25000
max	199.000000	130726.914100	175247.539500	1.163728	849.000000	239.00000

```
In [5]: goldman_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   TransactionID    800 non-null    int64  
 1   CustomerID       800 non-null    object  
 2   AccountID        800 non-null    object  
 3   AccountType      800 non-null    object  
 4   TransactionType  800 non-null    object  
 5   Product          800 non-null    object  
 6   Firm             800 non-null    object  
 7   Region           800 non-null    object  
 8   Manager          800 non-null    object  
 9   TransactionDate  800 non-null    object  
 10  TransactionAmount 800 non-null    float64 
 11  AccountBalance   800 non-null    float64 
 12  RiskScore        800 non-null    float64 
 13  CreditRating     800 non-null    int64  
 14  TenureMonths     800 non-null    int64  
dtypes: float64(3), int64(3), object(9)
memory usage: 93.9+ KB
```

```
In [6]: goldman_df.isnull().sum()
```

```
Out[6]: TransactionID      0
CustomerID        0
AccountID         0
AccountType       0
TransactionType   0
Product          0
Firm             0
Region           0
Manager          0
TransactionDate  0
TransactionAmount 0
AccountBalance   0
RiskScore         0
CreditRating      0
TenureMonths     0
dtype: int64
```

```
In [7]: for col in goldman_df.columns:
    if goldman_df[col].nunique()<10:
        print(goldman_df[col].value_counts())
        print("-"*20)
```

```

AccountType
Loan      218
Credit     206
Savings    204
Current    172
Name: count, dtype: int64
-----
TransactionType
Withdrawal 207
Payment    200
Deposit    199
Transfer   194
Name: count, dtype: int64
-----
Product
Home Loan 176
Credit Card 168
Mutual Fund 162
Personal Loan 153
Savings Account 141
Name: count, dtype: int64
-----
Firm
Firm B    168
Firm C    165
Firm D    164
Firm E    153
Firm A    150
Name: count, dtype: int64
-----
Region
West      176
East      162
South     158
North     156
Central   148
Name: count, dtype: int64
-----
Manager
Manager 1 218
Manager 3 209
Manager 2 193
Manager 4 180
Name: count, dtype: int64
-----
```

In [8]: `goldman_df.duplicated().sum()`

Out[8]: 0

In [9]: `# If you want to drop duplicates sum:
df = goldman_df.drop_duplicates()
df.shape`

Out[9]: (800, 15)

Task 1: Data Cleaning and Formatting

In [10]: `# 1.1 Remove/treat any special characters or non-numeric entries from financial fields.`

```

goldman_df["TransactionAmount"] = goldman_df["TransactionAmount"].abs()
financial_cols = ["TransactionAmount", "AccountBalance"]
for col in financial_cols:
    goldman_df[col] = pd.to_numeric(goldman_df[col], errors="coerce")
goldman_df.columns
goldman_df.head()
```

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Product	Firm	Region	Manager	TransactionDate	TransactionA
0	33	CUST6549	ACC1234	Credit	Withdrawal	Savings Account	Firm C	Central	Manager 1	21-10-2023	87481
1	177	CUST2942	ACC52650	Credit	Withdrawal	Home Loan	Firm A	East	Manager 3	20-06-2023	20311
2	178	CUST6776	ACC45101	Current	Deposit	Personal Loan	Firm C	South	Manager 3	02-01-2023	10481
3	173	CUST2539	ACC88252	Current	Withdrawal	Mutual Fund	Firm A	Central	Manager 2	25-07-2023	45121
4	67	CUST2626	ACC21878	Savings	Withdrawal	Home Loan	Firm C	Central	Manager 4	25-07-2023	42361



```
In [11]: # 1.2 Convert currency amounts into numerical format.

goldman_df.columns = goldman_df.columns.str.strip()

text_cols = goldman_df.select_dtypes(include="object").columns
goldman_df[text_cols] = goldman_df[text_cols].apply(lambda s: s.str.strip())
goldman_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   TransactionID    800 non-null    int64  
 1   CustomerID       800 non-null    object  
 2   AccountID        800 non-null    object  
 3   AccountType      800 non-null    object  
 4   TransactionType  800 non-null    object  
 5   Product          800 non-null    object  
 6   Firm             800 non-null    object  
 7   Region           800 non-null    object  
 8   Manager          800 non-null    object  
 9   TransactionDate  800 non-null    object  
 10  TransactionAmount 800 non-null    float64 
 11  AccountBalance   800 non-null    float64 
 12  RiskScore        800 non-null    float64 
 13  CreditRating     800 non-null    int64  
 14  TenureMonths     800 non-null    int64  
dtypes: float64(3), int64(3), object(9)
memory usage: 93.9+ KB
```

```
In [12]: # 1.3 Validate and format date columns.

goldman_df["TransactionDate"] = goldman_df["TransactionDate"].apply(pd.to_datetime, format="%d-%m-%Y")
goldman_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 800 entries, 0 to 799
Data columns (total 15 columns):
 #   Column           Non-Null Count  Dtype  
--- 
 0   TransactionID    800 non-null    int64  
 1   CustomerID       800 non-null    object  
 2   AccountID        800 non-null    object  
 3   AccountType      800 non-null    object  
 4   TransactionType  800 non-null    object  
 5   Product          800 non-null    object  
 6   Firm             800 non-null    object  
 7   Region           800 non-null    object  
 8   Manager          800 non-null    object  
 9   TransactionDate  800 non-null    datetime64[ns]
 10  TransactionAmount 800 non-null    float64 
 11  AccountBalance   800 non-null    float64 
 12  RiskScore        800 non-null    float64 
 13  CreditRating     800 non-null    int64  
 14  TenureMonths     800 non-null    int64  
dtypes: datetime64[ns](1), float64(3), int64(3), object(8)
memory usage: 93.9+ KB
```

```
In [13]: # 1.4 Ensure account types and transaction categories are standardized.

goldman_df["AccountType"] = goldman_df["AccountType"].str.strip().str.title()
goldman_df["TransactionType"] = goldman_df["TransactionType"].str.strip().str.title()
df[["TransactionType", "AccountType"]]
```

Out[13]:

	TransactionType	AccountType
0	Withdrawal	Credit
1	Withdrawal	Credit
2	Deposit	Current
3	Withdrawal	Current
4	Withdrawal	Savings
...
795	Withdrawal	Current
796	Withdrawal	Credit
797	Payment	Current
798	Payment	Current
799	Deposit	Loan

800 rows × 2 columns

Task 2: Descriptive Transactional Analysis

In [46]: *# 2.1 Calculate monthly and yearly summaries of total credits, debits, and net transaction volume.*

```
Credit = {"deposit"}
Debit = {"withdrawal", "transfer", "payment"}

goldman_df["Credit"] = np.where(
    goldman_df["TransactionType"].str.lower().isin(Credit),
    goldman_df["TransactionAmount"],
    0)

goldman_df["Debit"] = np.where(
    goldman_df["TransactionType"].str.lower().isin(Debit),
    goldman_df["TransactionAmount"],
    0)

goldman_df[["TransactionType", "TransactionAmount", "Credit", "Debit"]].head()
```

Out[46]:

	TransactionType	TransactionAmount	Credit	Debit
0	Withdrawal	87480.05448	0.00000	87480.05448
1	Withdrawal	20315.74505	0.00000	20315.74505
2	Deposit	10484.57165	10484.57165	0.00000
3	Withdrawal	45122.27373	0.00000	45122.27373
4	Withdrawal	42360.79878	0.00000	42360.79878

In [47]: *# Define Transaction Sign Function*

```
def txn_sign(txn_type: str) -> int:
    t = str(txn_type).lower()

    if any(k in t for k in Debit):
        return -1
    if any(k in t for k in Credit):
        return 1
    return 1

goldman_df["Sign"] = goldman_df["TransactionType"].apply(txn_sign)
goldman_df["Net_Amount"] = goldman_df["TransactionAmount"] * goldman_df["Sign"]

# Extract Month and Year
goldman_df["Year"] = goldman_df["TransactionDate"].dt.year
goldman_df["Month"] = goldman_df["TransactionDate"].dt.to_period("M").astype(str)

Monthly_Summary = goldman_df.groupby("Month").agg(
    TotalCredits = ("TransactionAmount", lambda s: s[goldman_df.loc[s.index, "Sign"] == 1].sum()),
    TotalDebits = ("TransactionAmount", lambda s: s[goldman_df.loc[s.index, "Sign"] == -1].sum()),
    Net = ("Net_Amount", "sum")).reset_index()

Yearly_Summary = goldman_df.groupby("Year").agg(
    TotalCredits = ("TransactionAmount", lambda s: s[goldman_df.loc[s.index, "Sign"] == 1].sum()),
    TotalDebits = ("TransactionAmount", lambda s: s[goldman_df.loc[s.index, "Sign"] == -1].sum()),
    Net = ("Net_Amount", "sum")).reset_index()

goldman_df.head()
```

Out[47]:

	TransactionID	CustomerID	AccountID	AccountType	TransactionType	Product	Firm	Region	Manager	TransactionDate	...	De
0	33	CUST6549	ACC12334	Credit	Withdrawal	Savings Account	Firm C	Central	Manager 1	2023-10-21	...	87480.05
1	177	CUST2942	ACC52650	Credit	Withdrawal	Home Loan	Firm A	East	Manager 3	2023-06-20	...	20315.74
2	178	CUST6776	ACC45101	Current	Deposit	Personal Loan	Firm C	South	Manager 3	2023-01-02	...	0.00
3	173	CUST2539	ACC88252	Current	Withdrawal	Mutual Fund	Firm A	Central	Manager 2	2023-07-25	...	45122.27
4	67	CUST2626	ACC21878	Savings	Withdrawal	Home Loan	Firm C	Central	Manager 4	2023-07-25	...	42360.79

5 rows × 26 columns

In [39]: # 2.2 Plot trends in total credits vs. debits over time.

```


sns.set_style(style="whitegrid")
plt.figure(figsize=(10,5))

# Plot Lines
sns.lineplot(
    data=Monthly_Summary,
    x="Month",
    y="TotalCredits",
    marker="o",
    label="Total Credits",
    color= "Orange")

sns.lineplot(
    data=Monthly_Summary,
    x="Month",
    y="TotalDebits",
    marker="o",
    label="Total Debits",
    color = "Brown")

plt.title("Total Credits vs Debit Monthly Trend")
plt.xlabel("Month")
plt.ylabel("Transaction Amount")
plt.xticks(rotation=45)
plt.legend()
plt.tight_layout()
plt.show()

```



In [49]: # 2.3 Identify top and bottom performing accounts based on net inflow.

```

# Top performing accounts

Top_Accounts= goldman_df.sort_values(by="Net_Amount",ascending=False).head(5)
print(f"Top Performing Accounts:\n{Top_Accounts[['Net_Amount','AccountBalance']]}")


# Bottom Performing Accounts

```

```
Bottom_Accounts = goldman_df.sort_values(by="Net_Amount", ascending=True).head(5)
print(f"Bottom Performing Accounts:\n{Bottom_Accounts[['Net_Amount', 'AccountBalance']]})")
```

Top Performing Accounts:

	Net_Amount	AccountBalance
741	130149.3799	51092.98942
237	127718.9013	33532.61157
234	122856.4830	93754.78157
379	118019.7086	42088.94404
269	113221.7353	32952.55399

	Net_Amount	AccountBalance
790	-130726.9141	116386.890400
80	-130475.1663	14174.176850
338	-129983.6051	42085.195500
319	-128816.2360	150199.078600
320	-127907.5731	9228.815058

In [18]: # 2.4 Identify and flag accounts as dormant or inactive if there is a gap of two months or more between consecutive transactions

```
goldman_df_sorted= goldman_df.sort_values(['AccountID','TransactionDate'])
goldman_df_sorted["PrevDate"] = goldman_df_sorted.groupby('AccountID')['TransactionDate'].shift(1)

goldman_df_sorted['GapDays'] = (
    goldman_df_sorted['TransactionDate'] - goldman_df_sorted['PrevDate']).dt.days

goldman_df_sorted["DormantFlag"] = goldman_df_sorted["GapDays"] >= 60 #If gap days is more than 60(2 months) its considered as dormant

Dormant_accounts = (goldman_df_sorted.groupby("AccountID")["DormantFlag"].
                     any().reset_index(name="Is-Dormant"))

goldman_df.head()
print(Dormant_accounts)
```

	AccountID	Is-Dormant
0	ACC10117	True
1	ACC10996	True
2	ACC11062	True
3	ACC11188	True
4	ACC11285	False
..
189	ACC97225	True
190	ACC97411	True
191	ACC99117	True
192	ACC99409	True
193	ACC99549	True

[194 rows x 2 columns]

Task 3: Customer Profile Building

In [19]: # 3.1 Group accounts by activity levels: High, Medium, Low based on transaction frequency on your analysis and rubrics. Do not forget to mention the rubric in the headings.

```
Transaction_frequency = (goldman_df.groupby(["CustomerID","Month"]).
                           size().reset_index(name="TransactionCount"))
print(Transaction_frequency)

Avg_Monthly_Transaction= (Transaction_frequency.groupby("CustomerID")["TransactionCount"].
                           mean().reset_index(name="Avg_Monthly_Transaction"))
print(Avg_Monthly_Transaction)

def ActivityLevel(x):
    if x >= 20:
        return "High"
    if x >= 5:
        return "Medium"
    else:
        return "Low"

Avg_Monthly_Transaction["ActivityLevel"] = Avg_Monthly_Transaction["Avg_Monthly_Transaction"].apply(ActivityLevel)

print((Avg_Monthly_Transaction["ActivityLevel"]).head(100))
```

```
CustomerID Month TransactionCount
0 CUST1042 2023-03 1
1 CUST1042 2023-04 1
2 CUST1042 2023-05 2
3 CUST1042 2023-07 1
4 CUST1042 2023-10 1
.. ...
700 CUST9843 2024-01 1
701 CUST9843 2024-02 1
702 CUST9962 2023-09 2
703 CUST9962 2024-01 1
704 CUST9962 2024-04 2
```

[705 rows x 3 columns]

	CustomerID	Avg_Monthly_Transaction
0	CUST1042	1.142857
1	CUST1114	1.333333
2	CUST1121	1.333333
3	CUST1189	1.000000
4	CUST1223	1.000000
..
183	CUST9683	1.750000
184	CUST9731	1.500000
185	CUST9754	1.333333
186	CUST9843	1.000000
187	CUST9962	1.666667

[188 rows x 2 columns]

0	Low
1	Low
2	Low
3	Low
4	Low
..	...
95	Low
96	Low
97	Low
98	Low
99	Low

Name: ActivityLevel, Length: 100, dtype: object

```
In [48]: # 3.2 Segment customers by average balance and transaction volume.
Account_metrics = (goldman_df.groupby("AccountID").agg(
    Total_Credit=("Credit", "sum"),
    Total_Debit=("Debit", "sum"),
    Transaction_Volume=("TransactionAmount", "sum")).

assign(Net_Inflow = lambda x: x["Total_Credit"] - x["Total_Debit"]).reset_index()

Account_metrics.head()
```

```
Out[48]:
```

	AccountID	Total_Credit	Total_Debit	Transaction_Volume	Net_Inflow
0	ACC10117	142170.20378	57310.763650	199480.967430	84859.440130
1	ACC10996	62580.86356	188158.687390	250739.550950	-125577.823830
2	ACC11062	0.00000	27189.136160	27189.136160	-27189.136160
3	ACC11188	45748.34156	211828.262030	257576.603590	-166079.920470
4	ACC11285	0.00000	96729.609841	96729.609841	-96729.609841

```
In [50]: # 3.3 Create profiles for:
# a) High-net inflow accounts:

High_inflow = Account_metrics["Net_Inflow"].quantile(0.75)
Low_inflow = Account_metrics["Net_Inflow"].quantile(0.25)

High_volume = Account_metrics["Transaction_Volume"].quantile(0.75)
Low_volume = Account_metrics["Transaction_Volume"].quantile(0.25)

def customer_seg(row):
    if row["Net_Inflow"] >= High_inflow or row["Transaction_Volume"] >= High_volume:
        return "High Value"
    elif row["Net_Inflow"] <= Low_inflow or row["Transaction_Volume"] <= Low_volume:
        return "Low Value"
    else:
        return "Medium Value"

Account_metrics["Customer_segment"] = Account_metrics.apply(customer_seg, axis=1)
Account_metrics.head()
```

Out[50]:	AccountID	Total_Credit	Total_Debit	Transaction_Volume	Net_Inflow	Customer_segment
0	ACC10117	142170.20378	57310.763650	199480.967430	84859.440130	High Value
1	ACC10996	62580.86356	188158.687390	250739.550950	-125577.823830	Medium Value
2	ACC11062	0.00000	27189.136160	27189.136160	-27189.136160	Low Value
3	ACC11188	45748.34156	211828.262030	257576.603590	-166079.920470	Medium Value
4	ACC11285	0.00000	96729.609841	96729.609841	-96729.609841	Low Value

```
In [54]: Account_metrics = (goldman_df.groupby("AccountID").agg(avg_balance=("AccountBalance", "mean"),
                                                               transaction_volume=("TransactionAmount", "sum")).reset_index())
Account_metrics.head()
```

Out[54]:	AccountID	avg_balance	transaction_volume
0	ACC10117	70107.007957	199480.967430
1	ACC10996	43568.008084	250739.550950
2	ACC11062	38137.132610	27189.136160
3	ACC11188	69652.151044	257576.603590
4	ACC11285	97401.348560	96729.609841

In [23]:

Out[23]:	CustomerID	AvgBalance	Total_Volume	Total_Net	Trans_Count
0	CUST1042	80435.167530	432345.676360	-284936.864820	8
1	CUST1114	60127.893845	385802.565919	-268029.670739	8
2	CUST1121	83213.472732	398205.117510	-398205.117510	4
3	CUST1189	63648.890237	117339.418890	-117339.418890	3
4	CUST1223	88186.197605	67776.465800	29853.686720	2

```
In [24]: High_net_inflow = Customer_Profile.sort_values("Total_Net", ascending=False).head(10)
print(High_net_inflow)
```

	CustomerID	AvgBalance	Total_Volume	Total_Net	Trans_Count
34	CUST2457	77828.604797	172045.102170	172045.102170	3
92	CUST5121	25888.608522	283940.122520	165814.500740	4
61	CUST3378	72427.257926	355094.057637	149915.581643	8
50	CUST3006	67652.686910	264585.713400	107755.139840	4
79	CUST4461	59922.191420	198753.919640	105440.304080	3
90	CUST4884	118249.524200	88990.271030	88990.271030	1
19	CUST1776	64601.246775	250310.765680	61632.060360	4
52	CUST3041	65911.485433	329713.516090	58656.963290	6
168	CUST9038	76172.026766	295900.217580	56552.403640	5
113	CUST6029	57679.042213	110182.364544	54734.730336	3

In [25]: # b) High-frequency Low-balance accounts:

```
Customer_Profile["Balance_Bucket"] = pd.qcut(Customer_Profile["AvgBalance"], 3, labels=["Low-Bal", "Mid-Bal", "High-Bal"])
Customer_Profile["Volume_Bucket"] = pd.qcut(Customer_Profile["Total_Volume"], 3, labels=["Low-Volume", "Mid-Volume", "High-Volume"])
print(Customer_Profile["Balance_Bucket"], Customer_Profile["Volume_Bucket"])

Customer_Profile.head()
```

```

0      High-Bal
1      Low-Bal
2      High-Bal
3      Low-Bal
4      High-Bal
...
183     Mid-Bal
184     High-Bal
185     High-Bal
186     Mid-Bal
187     High-Bal
Name: Balance_Bucket, Length: 188, dtype: category
Categories (3, object): ['Low-Bal' < 'Mid-Bal' < 'High-Bal'] 0      High-Volume
1      High-Volume
2      High-Volume
3      Low-Volume
4      Low-Volume
...
183     High-Volume
184     High-Volume
185     Mid-Volume
186     Mid-Volume
187     High-Volume
Name: Volume_Bucket, Length: 188, dtype: category
Categories (3, object): ['Low-Volume' < 'Mid-Volume' < 'High-Volume']

```

Out[25]:

	CustomerID	AvgBalance	Total_Volume	Total_Net	Trans_Count	Balance_Bucket	Volume_Bucket
0	CUST1042	80435.167530	432345.676360	-284936.864820	8	High-Bal	High-Volume
1	CUST1114	60127.893845	385802.565919	-268029.670739	8	Low-Bal	High-Volume
2	CUST1121	83213.472732	398205.117510	-398205.117510	4	High-Bal	High-Volume
3	CUST1189	63648.890237	117339.418890	-117339.418890	3	Low-Bal	Low-Volume
4	CUST1223	88186.197605	67776.465800	29853.686720	2	High-Bal	Low-Volume

In [26]:

```
High_freq_low_bal = Customer_Profile[(Customer_Profile["Trans_Count"] >= Customer_Profile["Trans_Count"].quantile(0.75)) &
                                      (Customer_Profile["AvgBalance"] <= Customer_Profile["AvgBalance"].quantile(0.25))]
print(High_freq_low_bal)
```

```

CustomerID      AvgBalance      Total_Volume      Total_Net      Trans_Count \
1      CUST1114    60127.893845    385802.565919    -268029.670739          8
53     CUST3059    59259.390588    296413.616370    -296413.616370          6
54     CUST3069    48347.423620    179440.309010    -97487.457390          5
80     CUST4584    52584.244758    298235.343843    -64988.794803          5
94     CUST5253    47802.857312    177843.520130    -157496.948350          5
99     CUST5428    58148.843446    261052.014510    -212246.224430          8
101    CUST5545    61151.359039    265333.196444    -265333.196444          7
109    CUST5912    61966.849233    354409.131240    -172474.809920          6
111    CUST5920    61623.184390    554945.597570    -338546.753770          7
130    CUST6937    52171.166596    220399.339887    -36247.379153          5
153    CUST8155    49397.810737    266000.319070    -139503.197730          5
154    CUST8250    51282.494454    243326.951520    -158946.685580          5
157    CUST8288    59954.562314    346459.485800    -275515.697980          5
162    CUST8461    48295.398164    336834.753390    -336834.753390          5

Balance_Bucket  Volume_Bucket
1      Low-Bal    High-Volume
53     Low-Bal    High-Volume
54     Low-Bal    Mid-Volume
80     Low-Bal    High-Volume
94     Low-Bal    Mid-Volume
99     Low-Bal    Mid-Volume
101    Low-Bal    High-Volume
109    Low-Bal    High-Volume
111    Low-Bal    High-Volume
130    Low-Bal    Mid-Volume
153    Low-Bal    High-Volume
154    Low-Bal    Mid-Volume
157    Low-Bal    High-Volume
162    Low-Bal    High-Volume

```

In [27]:

```
#c) Accounts with negative or near-zero balances:
Neg_or_near_zero_bal= Customer_Profile[(Customer_Profile["AvgBalance"] <= 0)]
print(Neg_or_near_zero_bal)
```

```
Empty DataFrame
Columns: [CustomerID, AvgBalance, Total_Volume, Total_Net, Trans_Count, Balance_Bucket, Volume_Bucket]
Index: []
```

Task 4: Financial Risk Identification

In [28]:

```
# 4.1 Track accounts with frequent Large withdrawals or overdrafts.
```

```
# Filter withdrawal transactions
Withdrawals = goldman_df[goldman_df["TransactionType"].str.contains("withdraw", na=False)]

# Define large withdrawal threshold
Large_withdrawal_threshold = Withdrawals["TransactionAmount"].quantile(0.90)

# Flag large withdrawals
Withdrawals["LargeWithdrawalFlag"] = (Withdrawals["TransactionAmount"] >= Large_withdrawal_threshold)

Large_withdrawals = (Withdrawals[Withdrawals["LargeWithdrawalFlag"]].groupby("AccountID")
                     .size().reset_index(name="LargeWithdrawalCount"))

print(Large_withdrawals)
```

Empty DataFrame
Columns: [AccountID, LargeWithdrawalCount]
Index: []

In [29]: # Overdraft- Accounts with negative balance

```
Overdraft_accts = goldman_df[goldman_df["AccountBalance"] < 0].groupby("AccountID").size().reset_index(name="OverdraftTransCount")
print(Overdraft_accts)
```

	AccountID	OverdraftTransCount
0	ACC16241	1
1	ACC19178	1
2	ACC23736	1
3	ACC26973	1
4	ACC28154	1
5	ACC28292	2
6	ACC29477	1
7	ACC33287	1
8	ACC49774	1
9	ACC58667	1
10	ACC70314	1
11	ACC77533	1
12	ACC83005	1
13	ACC88449	1
14	ACC94242	1

In [30]: #4.2 Calculate balance volatility using standard deviation or coefficient of variation.

```
Bal_volt = goldman_df.groupby("AccountID")["AccountBalance"].agg(["mean", "std"]).reset_index()
Bal_volt["CV"] = Bal_volt["std"] / Bal_volt["mean"].replace(0, np.nan)
print(Bal_volt)
```

	AccountID	mean	std	CV
0	ACC10117	70107.007957	25886.972758	0.369249
1	ACC10996	43568.008084	9434.002316	0.216535
2	ACC11062	38137.132610	3208.737888	0.084137
3	ACC11188	69652.151044	35494.660810	0.509599
4	ACC11285	97401.348560	55922.732441	0.574147
..
189	ACC97225	38652.306677	28069.592780	0.726207
190	ACC97411	55978.315635	7871.678922	0.140620
191	ACC99117	47228.185087	20780.582578	0.440004
192	ACC99409	83743.915565	21429.756821	0.255896
193	ACC99549	68641.201433	26251.797058	0.382450

[194 rows x 4 columns]

In [31]: # 4.3 Use IQR or z-score methods to detect anomalies.

```
# z-score :
import scipy.stats as stats

goldman_df["Z_Score"] = np.abs(stats.zscore(goldman_df['TransactionAmount']))
z_anomalies = goldman_df[goldman_df['Z_Score'] > 3]

print(goldman_df['Z_Score'])

goldman_df["Account_Z_Score"] = np.abs(stats.zscore(goldman_df['AccountBalance']))
z_anomalies = goldman_df[goldman_df['Account_Z_Score'] > 3]

print(goldman_df['Account_Z_Score'])
```

```

0      1.266887
1      1.148219
2      1.501730
3      0.256222
4      0.355519
...
795     2.431868
796     1.642768
797     0.389231
798     0.969360
799     1.018040
Name: Z_Score, Length: 800, dtype: float64
0      0.048835
1      1.457611
2      0.870504
3      1.228560
4      1.600137
...
795     0.008947
796     0.762771
797     1.094223
798     0.926659
799     0.368842
Name: Account_Z_Score, Length: 800, dtype: float64

```

In [32]: # 4.4 Highlight customers with irregular or suspicious transaction behavior.

```
Customer_stats = (goldman_df.groupby("CustomerID")["TransactionAmount"].agg(["mean", "std"]).reset_index())
print(Customer_stats)
```

	CustomerID	mean	std
0	CUST1042	54043.209545	30565.484684
1	CUST1114	48225.320740	26398.893405
2	CUST1121	99551.279378	6457.060342
3	CUST1189	39113.139630	14202.571329
4	CUST1223	33888.232900	21109.744323
..
183	CUST9683	38898.896979	21953.451838
184	CUST9731	62775.094836	39174.186502
185	CUST9754	23692.643309	13341.603086
186	CUST9843	50033.957967	17746.697147
187	CUST9962	57530.183012	39028.819945

[188 rows x 3 columns]

In [33]: goldman_df = goldman_df.drop(columns=["mean", "std"], errors="ignore")
goldman_df = goldman_df.merge(Customer_stats, on="CustomerID", how="left")
goldman_df.head()

```
goldman_df["Suspicious_Flag"] = goldman_df.apply(
    lambda row: "Suspicious"
    if pd.notna(row["std"]) and
       row["TransactionAmount"] > row["mean"] + 3 * row["std"]
    else "Normal",
    axis=1)
Suspicious_Customers = goldman_df[goldman_df["Suspicious_Flag"] == "Suspicious"][["CustomerID", "TransactionAmount", "mean", "std"]]
print(Suspicious_Customers)
```

Empty DataFrame

Columns: [CustomerID, TransactionAmount, mean, std]

Index: []

Task 5: Visualisation


In [34]: #5.1 Conduct extensive exploratory data analysis with attractive visualizations for your findings.

```
# 1) Histogram for Account Balance
sns.set_style("darkgrid")
plt.figure(figsize=(8, 5))

ax = sns.histplot(
    df["AccountBalance"],
    bins=10,
    kde=True)

for patch in ax.patches:
    height = patch.get_height()
    if height > 0:
        ax.annotate(
            f'{int(height)}',
            (patch.get_x() + patch.get_width() / 2, height),
            ha='center',
            va='bottom',
            fontsize=10,
            color='black')
plt.title("AccountBalance Distribution")
plt.xlabel("AccountBalance")
```

```
plt.ylabel("Frequency")
plt.tight_layout()
plt.show()
```




In [35]: # 2) Scatter Plot: Customer Transaction Volume vs Avg Balance

```
Customer_Profile["AvgBalance_Round"] = (
    Customer_Profile["AvgBalance"] // 10000 * 10000).astype(int)
sns.set_style("whitegrid")
plt.figure(figsize=(9,6))

sns.scatterplot(
    data=Customer_Profile,
    x="Total_Volume",
    y="AvgBalance",
    hue= "AvgBalance_Round",
    s=30,
    alpha=0.75,
    palette="colorblind")

plt.title("Customer Transaction Volume vs Avg Balance")
plt.xlabel("Total Transaction Volume")
plt.ylabel("Avg Balance")
plt.grid(True)
plt.show()
```



In [36]: # 3) Boxplot: Transaction Amount by Activity Level

```
merged = goldman_df.merge(
    Avg_Monthly_Transaction[["CustomerID", "ActivityLevel"]],
    on="CustomerID",
    how="left")


plt.figure(figsize=(8,5))

# Draw the boxplot
sns.boxplot(
    data=merged,
    x="ActivityLevel",
    y="TransactionAmount",
    showfliers=False )

# Calculate min and max for each ActivityLevel
grouped = merged.groupby("ActivityLevel")["TransactionAmount"]
mins = grouped.min().values
maxs = grouped.max().values
positions = range(len(mins))

# Annotate min and max
for pos, min_val, max_val in zip(positions, mins, maxs):
    plt.text(pos, min_val, f'Min: {min_val:.2f}', horizontalalignment='center', color='blue')
    plt.text(pos, max_val, f'Max: {max_val:.2f}', horizontalalignment='center', color='green')

plt.title("Transaction Amount by Activity Level")
plt.xlabel("Activity Level")
plt.ylabel("Transaction Amount")
plt.show()
```




In [37]: # 4) Countplot to visualize Dormant vs Active customers

```
# Map boolean to string for better labels
Dormant_accounts['Status'] = Dormant_accounts['Is-Dormant'].map({True: 'Dormant', False: 'Active'})

plt.figure(figsize=(6,4))
sns.countplot(data=Dormant_accounts, x='Status', hue='Status', palette='deep')
plt.title("Dormant vs Active Customers")
plt.xlabel("Customer Status")
plt.ylabel("Number of Customers")

# Add data labels
for p in plt.gca().patches:
    plt.gca().annotate(int(p.get_height()), (p.get_x() + p.get_width()/2., p.get_height()),
                      ha='center', va='bottom')

plt.show()
```



Task 6: Hypothesis Testing

```
In [55]: # 6.1 Test whether high-volume transaction accounts have statistically higher average balances than low-volume accounts.
```

```
from scipy.stats import ttest_ind

# Split groups
q75 = Customer_Profile["Total_Volume"].quantile(0.75)
q25 = Customer_Profile["Total_Volume"].quantile(0.25)

high_volume = Customer_Profile[Customer_Profile["Total_Volume"] >= q75]["AvgBalance"]
low_volume = Customer_Profile[Customer_Profile["Total_Volume"] <= q25]["AvgBalance"]

t_stat, p_value = ttest_ind(high_volume, low_volume,
    alternative="greater", # one-tailed
    equal_var=False)       # Welch's test

print("T-statistic:", t_stat)
print("P-value:", p_value)
alpha = 0.05
if p_value < alpha:
    print("Decision: Reject H0 → High-volume accounts have higher average balances.")
else:
    print("Decision: Fail to reject H0 → No significant difference found.")
```

T-statistic: -0.6458575272439284

P-value: 0.7395549106185177

Decision: Fail to reject H0 → No significant difference found.

```
In [56]: # 6.2 Conduct hypothesis testing based on segmentation.
```

```
q75 = Customer_Profile["Trans_Count"].quantile(0.75)
q25 = Customer_Profile["Trans_Count"].quantile(0.25)

goldman_df["Customer_Segment"] = Customer_Profile["Trans_Count"].apply(lambda x: "High Volume" if x >= q75 else "Low Volume")

high_volume_bal = goldman_df[goldman_df["Customer_Segment"] == "High Volume"]["AccountBalance"].dropna()
low_volume_bal = goldman_df[goldman_df["Customer_Segment"] == "Low Volume"]["AccountBalance"].dropna()

t_stat, p_value = ttest_ind(high_volume_bal, low_volume_bal, equal_var=False)
print(t_stat, p_value)

if p_value < 0.05: print("Decision: Reject H0: Significant difference exists")
else: print("Decision: Fail to reject H0: No significant difference")
```

0.16229630215021854 0.8712636910790821

Decision: Fail to reject H0: No significant difference

Findings and Recommendation:

Findings:

1. Customer Activity: Most customers are Medium activity; fewer are Low or High activity.
2. Transaction Amounts: High-activity accounts have the largest transaction ranges, Low-activity accounts have mostly small transactions.
3. Debit vs Credit Trends: High-activity customers conduct mostly debits, Low-activity accounts have fewer transactions and more credits.

4. High-Risk Accounts: Accounts with long transaction gaps or unusually high transaction spikes are flagged as risky.

Recommendations:

1. Engage Low-Activity Customers – Encourage them to use their accounts more regularly.
2. Monitor High-Risk Accounts – Track accounts with unusual activity to prevent issues.
3. Customize Services by Segment – Offer products based on customer behavior and account type.